

Compilation wrapper for meanfield_lowrank.tex

1 Low-rank channel mixing: where the factor r enters (and why Picard is unchanged)

(proofs for flagged results are deferred to the appendix / supplementary material).

This note isolates the only extra bookkeeping induced by the channel-mixing / low-rank structure

$$H_2(t, c_2; X, W) = \sum_{k=1}^r L_{c_2, k} m_k(t; X, W), \quad m_k(t; X, W) = \mathbb{E}_{C_1}[w_1(t, C_1, k) \varphi_1(\langle f_1(C_1), X \rangle)],$$

relative to the standard mean-field ODE well-posedness template (a priori ψ_2 bounds, good/bad event decomposition, and Picard iteration).

1.1 A deterministic r -factor for low-rank sums

It is convenient to record row norms of L :

$$\|L\|_{\infty, 1} \equiv \sup_{c_2 \in \Omega_2} \sum_{k=1}^r |L_{c_2, k}|, \quad \|L\|_{\infty, 2} \equiv \sup_{c_2 \in \Omega_2} \left(\sum_{k=1}^r L_{c_2, k}^2 \right)^{1/2}.$$

Under the entrywise bound $\sup_{c_2, k} |L_{c_2, k}| \leq K$ (as in Assumption ??),

$$\|L\|_{\infty, 1} \leq rK, \quad \|L\|_{\infty, 2} \leq \sqrt{r} K.$$

Lemma 1 (Moment and ψ_2 control for low-rank sums). *Let $(U_k)_{k=1}^r$ be real random variables on a common probability space and $(a_k)_{k=1}^r \in \mathbb{R}^r$ deterministic. Then for every $m \geq 1$,*

$$\mathbb{E} \left[\left| \sum_{k=1}^r a_k U_k \right|^m \right]^{1/m} \leq \sum_{k=1}^r |a_k| \mathbb{E}[|U_k|^m]^{1/m} \leq \left(\sum_{k=1}^r |a_k| \right) \max_{1 \leq k \leq r} \mathbb{E}[|U_k|^m]^{1/m}.$$

Consequently, any “ ψ_2 -type” seminorm defined by $\sup_{m \geq 1} m^{-1/2} \mathbb{E}[|\cdot|^m]^{1/m}$ satisfies

$$\sup_{m \geq 1} \frac{1}{\sqrt{m}} \mathbb{E} \left[\left| \sum_{k=1}^r a_k U_k \right|^m \right]^{1/m} \leq \left(\sum_{k=1}^r |a_k| \right) \max_{1 \leq k \leq r} \sup_{m \geq 1} \frac{1}{\sqrt{m}} \mathbb{E}[|U_k|^m]^{1/m}.$$

If one prefers an ℓ_2 version, then also $\sum_{k=1}^r |a_k| \leq \sqrt{r} (\sum_k a_k^2)^{1/2}$ yields a $\sqrt{r}$ factor.

1.2 Application to H_2 : $\|L\|_{\infty, 1}$ is the only new constant

Lemma 2 (Channel mixing only multiplies forward bounds by $\|L\|_{\infty, 1}$). *Assume φ_1 is bounded by K . Then for every $t \geq 0$ and every $c_2 \in \Omega_2$,*

$$\sup_{s \leq t} |H_2(s, c_2; X, W)| \leq \left(\sum_{k=1}^r |L_{c_2, k}| \right) \max_{1 \leq k \leq r} \sup_{s \leq t} |m_k(s; X, W)| \leq \|L\|_{\infty, 1} \max_{1 \leq k \leq r} \sup_{s \leq t} |m_k(s; X, W)|.$$

Moreover,

$$\sup_{s \leq t} |m_k(s; X, W)| \leq \mathbb{E}_{C_1} \left[\sup_{s \leq t} |w_1(s, C_1, k)| |\varphi_1(\langle f_1(C_1), X \rangle)| \right] \leq K \mathbb{E}_{C_1} \left[\sup_{s \leq t} |w_1(s, C_1, k)| \right].$$

Combining these inequalities and taking moments yields, for every $m \geq 1$,

$$\mathbb{E}_X \left[\sup_{s \leq t} \sup_{c_2} |H_2(s, c_2; X, W)|^m \right]^{1/m} \leq K \|L\|_{\infty,1} \max_{1 \leq k \leq r} \mathbb{E}_{C_1} \left[\sup_{s \leq t} |w_1(s, C_1, k)|^m \right]^{1/m}.$$

In particular, in the ψ_2 -type calibration of Section 3,

$$\sqrt{50} \sup_{m \geq 1} \frac{1}{\sqrt{m}} \mathbb{E}_X \left[\sup_{s \leq t} \sup_{c_2} |H_2(s, c_2; X, W)|^m \right]^{1/m} \leq K \|L\|_{\infty,1} \llbracket w_1 \rrbracket_{\psi,t}.$$

Remark 1 (The promised r factor). Under $\sup_{c_2,k} |L_{c_2,k}| \leq K$, Lemma ?? gives an explicit r -dependence

$$\llbracket H_2 \rrbracket_{\psi,t} \lesssim (rK) \llbracket w_1 \rrbracket_{\psi,t},$$

up to the same universal constants already present in the two-layer well-posedness proof. If one instead assumes a row ℓ_2 control $\|L\|_{\infty,2} \leq K$ (e.g. normalized rows), then the same steps yield a $\sqrt{r}$ scaling.

1.3 Where this plugs into the Picard / well-posedness template

In the Picard map F (Section 4), the mixing kernel L appears only through H_2 and through the backpropagated signal

$$B_k(s; X, W') = \mathbb{E}_{C_2} [L_{C_2,k} \varphi'_2(H_2(s, C_2; X, W')) w'_2(s, C_2)].$$

Because φ'_2 is bounded, B_k itself does not introduce any additional r factor (it has one fixed column k of L); the only place where r can enter is when controlling H_2 (and hence $\varphi_2(H_2)$ and $\varphi'_2(H_2)$ through Lipschitzness).

Concretely, in the two-layer roadmap of Section 5:

- Lemma C (Forward Lipschitz) uses a bound of the form

$$|H_2(W') - H_2(W'')| \leq \|L\|_{\infty,1} \max_k |m_k(W') - m_k(W'')| \leq K \|L\|_{\infty,1} \|w'_1 - w''_1\|_t,$$

so the constant in Lemma C is multiplied by at most $\|L\|_{\infty,1}$ (hence at most rK under entrywise control).

- Lemma D (Backward Lipschitz) inherits the same multiplicative factor when bounding drift differences involving $\varphi_2(H_2)$ or $\varphi'_2(H_2)$, because those differences are controlled by Lipschitz constants times $|H_2(W') - H_2(W'')|$ on the good event.

Therefore the difference inequality of the form

$$\|F(W') - F(W'')\|_t \leq C \int_0^t \left((1+B) \|W' - W''\|_s + e^{-cB^2} \right) ds$$

remains valid with the same structure, after replacing C by $C \text{ poly}(\|L\|_{\infty,1})$. The Picard iteration and the $B = \sqrt{m}$ choice are unchanged; only the polynomial function $K_0(t)$ (and hidden constants) now depend on $\|L\|_{\infty,1}$, i.e. on r through $\|L\|_{\infty,1} \leq r \sup_{c_2,k} |L_{c_2,k}|$.

2 Channel-wise partial functions via pushforward and disintegration

This section rewrites the r channel summaries $m_k(t; \cdot)$ as integrals against a measure μ_0 on the *untrained* first-layer features, together with a pushforward (or, equivalently, a backward kernel) that incorporates the *trained* first-layer weights w_1 .

2.1 Untrained features and their pushforward law

Let $w_0 : \Omega_1 \rightarrow \mathbb{R}^d$ denote the untrained first-layer map (in the main note, $w_0 = f_1$) and let ρ^1 be the law of C_1 . Define the pushforward measure on feature space

$$\mu_0 \equiv (w_0)_\# \rho^1 \quad \text{on } \mathbb{R}^d,$$

meaning that for any measurable $A \subseteq \mathbb{R}^d$, $\mu_0(A) = \rho^1(\{c_1 : w_0(c_1) \in A\})$.

2.2 The r partial functions as integrals over μ_0

For each channel $k \in \{1, \dots, r\}$, recall the partial function

$$m_k(t; X, W) = \mathbb{E}_{C_1} [w_1(t, C_1, k) \varphi_1(\langle w_0(C_1), X \rangle)].$$

It is often useful to view $m_k(t; \cdot)$ as an integral over μ_0 by packaging w_1 into a signed (or weighted) pushforward measure.

Definition 1 (Signed pushforward incorporating w_1). *For each $t \geq 0$ and each $k \in \{1, \dots, r\}$, define the (finite signed) measure $\mu_{1,k}^t$ on $\mathbb{R}^d$ by*

$$\mu_{1,k}^t \equiv (w_0)_\# (w_1(t, \cdot, k) \rho^1),$$

i.e. for any measurable $A \subseteq \mathbb{R}^d$,

$$\mu_{1,k}^t(A) = \int_{\Omega_1} \mathbb{1}\{w_0(c_1) \in A\} w_1(t, c_1, k) \rho^1(dc_1).$$

Lemma 3 (Integral representation over feature space). *For every $t \geq 0$, $k \in \{1, \dots, r\}$ and input X ,*

$$m_k(t; X, W) = \int_{\mathbb{R}^d} \varphi_1(\langle \theta, X \rangle) \mu_{1,k}^t(d\theta).$$

2.3 Backward kernel (disintegration) form

If one prefers to keep μ_0 fixed and represent only a *density* on top of it, one can use a disintegration of ρ^1 along the map w_0 . There exists a measurable Markov kernel $\kappa(dc_1 \mid \theta)$ on Ω_1 such that

$$\rho^1(dc_1) = \int_{\mathbb{R}^d} \kappa(dc_1 \mid \theta) \mu_0(d\theta), \quad \kappa(\{c_1 : w_0(c_1) = \theta\} \mid \theta) = 1 \quad \text{for } \mu_0\text{-a.e. } \theta.$$

Define the channel-wise conditional average of w_1 at feature location θ by

$$\bar{w}_{1,k}(t, \theta) \equiv \int_{\Omega_1} w_1(t, c_1, k) \kappa(dc_1 \mid \theta).$$

Then $\mu_{1,k}^t$ is absolutely continuous with respect to μ_0 with Radon–Nikodym derivative $\bar{w}_{1,k}(t, \cdot)$, in the sense that

$$\mu_{1,k}^t(d\theta) = \bar{w}_{1,k}(t, \theta) \mu_0(d\theta),$$

and hence the k th partial function can be written as the μ_0 -integral

$$m_k(t; X, W) = \int_{\mathbb{R}^d} \bar{w}_{1,k}(t, \theta) \varphi_1(\langle \theta, X \rangle) \mu_0(d\theta).$$

Remark 2 (Interpretation and special cases). *The family $(m_k(t; \cdot))_{k=1}^r$ can be seen as r scalar functions generated by the same untrained feature measure μ_0 and r channel-dependent coefficient fields $\bar{w}_{1,k}(t, \cdot)$ on feature space. If w_0 is injective, then $\kappa(dc_1 \mid \theta)$ is a Dirac mass at $c_1 = w_0^{-1}(\theta)$ and $\bar{w}_{1,k}(t, \theta) = w_1(t, w_0^{-1}(\theta), k)$.*

3 Time variation of the r partial functions and the induced PDE

Define for each channel $k \in \{1, \dots, r\}$ the partial function on input space

$$f_k(t, x) \equiv m_k(t; x, W) = \int_{\mathbb{R}^d} \bar{w}_{1,k}(t, \theta) \varphi_1(\langle \theta, x \rangle) \mu_0(d\theta).$$

We compute its time derivative and state the (measure-valued) evolution equation implied by the MF ODE for w_1 .

3.1 Evolution of the coefficient field $\bar{w}_{1,k}(t, \theta)$

From the MF ODE (cf. (??) in the main note), the first-layer weights satisfy

$$\partial_t w_1(t, c_1, k) = -\xi_1(t) \mathbb{E}_Z \left[d_L(Z; W(t)) \varphi_1(\langle w_0(c_1), X \rangle) \mathbb{E}_{C_2} \left[L_{C_2,k} \varphi_2' \left(H_2(t, C_2; X, W(t)) \right) w_2(t, C_2) \right] \right].$$

Introduce the channel- k backpropagated signal (a random variable through X)

$$B_k(t; X) \equiv \mathbb{E}_{C_2} \left[L_{C_2,k} \varphi_2' \left(H_2(t, C_2; X, W(t)) \right) w_2(t, C_2) \right].$$

Since the drift depends on c_1 only through $\theta = w_0(c_1)$, the conditional average $\bar{w}_{1,k}(t, \theta)$ evolves pointwise as

$$\partial_t \bar{w}_{1,k}(t, \theta) = -\xi_1(t) \mathbb{E}_{Z=(X,Y)} \left[d_L(Z; W(t)) \varphi_1(\langle \theta, X \rangle) B_k(t; X) \right]. \quad (1)$$

3.2 A measure-valued PDE for $\mu_{1,k}^t$

Recall $\mu_{1,k}^t(d\theta) = \bar{w}_{1,k}(t, \theta) \mu_0(d\theta)$. Equation (??) is equivalently the following linear “reaction” PDE for the signed measure $\mu_{1,k}^t$:

$$\partial_t \mu_{1,k}^t(d\theta) = -\xi_1(t) \mathbb{E}_{Z=(X,Y)} \left[d_L(Z; W(t)) \varphi_1(\langle \theta, X \rangle) B_k(t; X) \right] \mu_0(d\theta). \quad (2)$$

In weak form, for every bounded measurable test function $\psi : \mathbb{R}^d \rightarrow \mathbb{R}$,

$$\frac{d}{dt} \int_{\mathbb{R}^d} \psi(\theta) \mu_{1,k}^t(d\theta) = -\xi_1(t) \mathbb{E}_{Z=(X,Y)} \left[d_L(Z; W(t)) B_k(t; X) \int_{\mathbb{R}^d} \psi(\theta) \varphi_1(\langle \theta, X \rangle) \mu_0(d\theta) \right].$$

Because w_0 is frozen, there is no transport term (no $\nabla_\theta \cdot (\cdot)$): the feature locations stay fixed and only the signed density evolves.

3.3 Variation of $f_k(t, x)$ and an induced integro-differential equation

Differentiating under the integral and using (??) yields

$$\begin{aligned}\partial_t f_k(t, x) &= \int_{\mathbb{R}^d} \partial_t \bar{w}_{1,k}(t, \theta) \varphi_1(\langle \theta, x \rangle) \mu_0(d\theta) \\ &= -\xi_1(t) \mathbb{E}_{Z=(X,Y)} \left[d_L(Z; W(t)) B_k(t; X) \int_{\mathbb{R}^d} \varphi_1(\langle \theta, X \rangle) \varphi_1(\langle \theta, x \rangle) \mu_0(d\theta) \right].\end{aligned}\quad (3)$$

Define the kernel induced by the untrained feature measure μ_0 :

$$K_{\mu_0}(x, x') \equiv \int_{\mathbb{R}^d} \varphi_1(\langle \theta, x \rangle) \varphi_1(\langle \theta, x' \rangle) \mu_0(d\theta).$$

Then (??) becomes the compact evolution equation

$$\partial_t f_k(t, x) = -\xi_1(t) \mathbb{E}_{Z=(X,Y)} [d_L(Z; W(t)) B_k(t; X) K_{\mu_0}(x, X)], \quad k = 1, \dots, r. \quad (4)$$

This is an r -coupled integro-differential system in t (and in x through the kernel), closed once one expresses $B_k(t; X)$ and $d_L(Z; W(t))$ in terms of the network output $\hat{y}(\cdot; W(t))$, which itself depends on $(f_j)_{j=1}^r$ through $H_2 = \sum_{j=1}^r L_{\cdot,j} f_j$.

3.4 ReLU top layer: sign-coherent gating and a reinforcement heuristic

Assume the second-layer activation is ReLU, $\varphi_2(u) = u_+$, so that $\varphi_2'(u) = \mathbb{1}\{u > 0\}$ a.e. Then the channel- k backpropagated signal becomes

$$B_k(t; X) = \mathbb{E}_{C_2} [L_{C_2,k} w_2(t, C_2) \mathbb{1}\{H_2(t, C_2; X, W(t)) > 0\}]. \quad (5)$$

In words: $B_k(t; X)$ is a *gated mixture* of the k th column of L , weighted by the current top-layer weights w_2 , with the gate determined by H_2 .

Combining (??) with (??) shows that, for each k ,

$$\partial_t f_k(t, x) = -\xi_1(t) \mathbb{E}_{Z=(X,Y)} [d_L(Z; W(t)) K_{\mu_0}(x, X) \mathbb{E}_{C_2} [L_{C_2,k} w_2(t, C_2) \mathbb{1}\{H_2(t, C_2; X, W(t)) > 0\}]].$$

Remark 3 (When does “large f_j makes itself larger” hold?). *The intuition that a comparatively large f_j amplifies its own future growth goes through a sign-coherence regime:*

- if $L_{c_2,k} \geq 0$ (or at least does not change sign too much) for the relevant c_2 ,
- if $w_2(t, c_2) \geq 0$ (or has a controlled sign pattern aligned with $L_{c_2,\cdot}$),
- if the loss derivative $d_L(Z; W(t))$ has a consistent sign on the region of inputs currently being learned (e.g. in a residual-fitting phase for square loss),

then increasing one partial function $f_j(t, \cdot)$ tends to increase H_2 on the subset of (C_2, X) where $L_{C_2,j} > 0$, which increases the gate $\mathbb{1}\{H_2 > 0\}$ and hence increases $B_j(t; X)$. Since $\partial_t f_j$ is proportional to $B_j(t; X)$ through (??), this creates a positive feedback loop in that region.

Without such sign restrictions, reinforcement is not guaranteed: w_2 and d_L can change sign, and $L_{c_2,k}$ may have mixed signs, so the same mechanism can produce cancellation or even competition between channels.

3.5 Fixed random-sign L : an “emergent” sign-coherence in a one-spike reduction

The sign-coherence bullets in Remark ?? are sufficient, but in practice $L_{c_2,k}$ can have mixed signs (e.g. random at initialization and frozen). In that case, one can still obtain an *effective* sign-coherence once the dynamics at a given input location is dominated by a single channel. We explain the mechanism in a simplified “single spike” reduction.

Single spike reduction. Assume the input distribution is concentrated at one point x_* (or that we are in a local-kernel regime where the dynamics at x_* decouple), and set

$$f_k(t) \equiv f_k(t, x_*), \quad d(t) \equiv d_L((x_*, y_*); W(t)), \quad L_{c_2} \equiv (L_{c_2,1}, \dots, L_{c_2,r}) \in \mathbb{R}^r.$$

With ReLU $\varphi_2(u) = u_+$, the gate is $\mathbb{1}\{\langle L_{c_2}, f(t) \rangle > 0\}$ and the two relevant MF equations become

$$\partial_t f_k(t) = -\xi_1(t) d(t) B_k(t), \tag{6}$$

$$B_k(t) = \mathbb{E}_{C_2}[L_{C_2,k} w_2(t, C_2) \mathbb{1}\{\langle L_{C_2}, f(t) \rangle > 0\}], \quad k = 1, \dots, r, \tag{7}$$

$$\partial_t w_2(t, c_2) = -\xi_2(t) d(t) (\langle L_{c_2}, f(t) \rangle)_+. \tag{8}$$

If $w_2(0, \cdot) \geq 0$ and $-d(t) \geq 0$ (for instance in an underfitting phase for square loss), then (??) implies $w_2(t, \cdot)$ remains nonnegative and increases more for those c_2 with larger positive $\langle L_{c_2}, f(t) \rangle$.

Why mixed signs in L may still reinforce a dominant channel. Suppose at some time t the vector $f(t)$ is (approximately) one-sparse,

$$f(t) \approx a(t) e_j, \quad a(t) \in \mathbb{R}, \quad j \in \{1, \dots, r\}.$$

Then $\langle L_{c_2}, f(t) \rangle \approx a(t) L_{c_2,j}$ and the gate becomes $\mathbb{1}\{a(t) L_{c_2,j} > 0\}$. If (heuristically, or in an early-time approximation) $w_2(t, C_2)$ is not strongly *negatively* correlated with the sign of $L_{C_2,j}$, then conditioning on the gate makes the j th coordinate have a *nonzero mean*:

$$\mathbb{E}[L_{C_2,j} \mathbb{1}\{a L_{C_2,j} > 0\}] = \text{sign}(a) \mathbb{E}[|L_{C_2,j}|] / 2,$$

whereas for $k \neq j$, under approximate independence and symmetry across coordinates of L_{C_2} ,

$$\mathbb{E}[L_{C_2,k} \mathbb{1}\{a L_{C_2,j} > 0\}] \approx 0.$$

Plugging this into (??) suggests that, even when L has mixed signs, the gate selects a half-space primarily determined by the dominant coordinate j , and this selection biases $B_j(t)$ to have the *same sign* as $a(t)$ while keeping $B_k(t)$ small for $k \neq j$. Equation (??) then amplifies the dominant channel at that spike location. Moreover, (??) further increases $w_2(t, c_2)$ on the gated set $\{c_2 : \langle L_{c_2}, f(t) \rangle > 0\}$, strengthening the conditional bias in (??).

A rigorous one-sparse statement. The preceding intuition can be made fully rigorous in the exact one-sparse regime, under a mild symmetry/independence assumption on the frozen random vector L_{C_2} and a sign-consistent residual.

Assumption 1 (Random mixing vector, symmetric and independent). *The random vector $L_{C_2} = (L_{C_2,1}, \dots, L_{C_2,r})$ has independent coordinates, each symmetric about 0: $L_{C_2,k} \stackrel{d}{=} -L_{C_2,k}$ and $\mathbb{E}[L_{C_2,k}^2] < \infty$ for all k .*

Lemma 4 (Half-space symmetry identities). *Let U be a real random variable with $U \stackrel{d}{=} -U$ and $\mathbb{E}[|U|] < \infty$. Then for any $a \in \mathbb{R} \setminus \{0\}$,*

$$\mathbb{E}[U \mathbb{1}\{aU > 0\}] = \text{sign}(a) \frac{1}{2} \mathbb{E}[|U|].$$

If moreover $\mathbb{E}[U^2] < \infty$, then

$$\mathbb{E}[U^2 \mathbb{1}\{aU > 0\}] = \frac{1}{2} \mathbb{E}[U^2].$$

Proof deferred to appendix.

Lemma 5 (One-sparse invariant manifold and emergent sign-coherence). *Assume the one-spike system (??)–(??), Assumption ??, and that the residual has fixed sign on $[0, T]$: $d(t) \leq 0$ for all $t \in [0, T]$. Assume a nonnegative constant initialization in the top layer, $w_2(0, c_2) \equiv w_2^0 \geq 0$. Fix $j \in \{1, \dots, r\}$ and assume a one-sparse spike at $t = 0$ with positive amplitude, $f(0) = a_0 e_j$ with $a_0 > 0$. Then for all $t \in [0, T]$:*

- $f(t) = a(t) e_j$ for some scalar $a(t) \geq a_0$;
- for $k \neq j$, $B_k(t) = 0$;
- $B_j(t) > 0$ and admits the explicit representation

$$B_j(t) = \frac{w_2^0}{2} \mathbb{E}[|L_{C_2, j}|] + \frac{\mathbb{E}[L_{C_2, j}^2]}{2} \int_0^t \xi_2(s) (-d(s)) a(s) ds. \quad (9)$$

Consequently $\partial_t a(t) = -\xi_1(t) d(t) B_j(t) \geq 0$ on $[0, T]$.

Proof deferred to appendix.

Remark 4 (Interpretation of Lemma ??). *Even if $L_{C_2, j}$ has mixed signs, once $f(t) = a(t) e_j$ with $a(t) > 0$ the gate is exactly $\mathbb{1}\{L_{C_2, j} > 0\}$ and $w_2(t, c_2)$ increases on that half-space via (??). Lemma ?? then implies a strictly positive bias in $B_j(t)$, while for $k \neq j$ the independence and symmetry enforce $B_k(t) = 0$.*

Takeaway. In this one-spike regime, random sign patterns do not necessarily destroy reinforcement; instead, once one channel dominates, the ReLU gate can create an emergent sign-coherence by conditioning on a half-space determined by that dominant coordinate. This provides a plausible route to winner-take-all specialization without requiring $L_{c_2, k} \geq 0$ for all k .

Remark 5 (Heuristic link to specialization and “spikes”). *In one-dimensional problems, the gate $\mathbb{1}\{H_2(t, C_2; x) > 0\}$ can partition the input axis into regions that activate different subsets of c_2 (hence different linear combinations of the r channels through L). If different channels k become dominant on different regions, (??) shows they are trained using different effective signed measures over X , yielding a self-reinforcing specialization in distinct “slots”. This is a plausible mechanism for learning localized high-amplitude structures (spikes) at different locations, driven by the gating nonlinearity rather than by a globally smooth kernel alone.*

4 Toward a proof of spike specialization: a workable program

(this section is a proof program / conjectural roadmap; statements are conditional and may require additional assumptions).

This section records a proof “démarche” for the empirically observed phenomenon: different channels $k \in \{1, \dots, r\}$ learn localized high-amplitude structures at different input locations (“spike specialization”). The goal is to isolate statements that are both meaningful and provable from the dynamics (??)–(??).

4.1 Step 0: formalize what “specialization” and “spike” mean

There are at least two mathematically tractable targets.

(A) Channel specialization (winner-take-all at each location). Define the channel shares

$$s_k(t, x) \equiv \frac{|f_k(t, x)|}{\sum_{j=1}^r |f_j(t, x)| + \varepsilon}, \quad k = 1, \dots, r,$$

with $\varepsilon > 0$ for stability. One aims to show that for many x there exists $k^*(x)$ such that $s_{k^*(x)}(t, x) \rightarrow 1$ and $s_k(t, x) \rightarrow 0$ for $k \neq k^*(x)$. This captures specialization without committing to a particular notion of spike.

(B) Spike formation (localization / high-frequency growth). One aims to show that for some channels k , the function $f_k(t, \cdot)$ becomes localized, e.g. $\|f_k(t, \cdot)\|_\infty$ grows while $\|f_k(t, \cdot)\|_{L^1}$ stays controlled, or that a suitable high-frequency energy increases. This typically requires additional structure on μ_0 and φ_1 (e.g. wavelet-like features).

4.2 Step 1: reduce to an empirical (finite-sample) dynamical system

Let $x_1, \dots, x_n$ be the training inputs and consider the finite vector process

$$F_k(t) \equiv (f_k(t, x_1), \dots, f_k(t, x_n)) \in \mathbb{R}^n, \quad k = 1, \dots, r.$$

If one replaces $\mathbb{E}_X[\cdot]$ by the empirical average $n^{-1} \sum_{i=1}^n (\cdot)|_{X=x_i}$, then (??) yields an ODE on $\mathbb{R}^{r \times n}$:

$$\partial_t f_k(t, x_i) = -\xi_1(t) \frac{1}{n} \sum_{j=1}^n d_L((x_j, y_j); W(t)) B_k(t; x_j) K_{\mu_0}(x_i, x_j).$$

This finite-dimensional setting is a natural starting point for a rigorous analysis of specialization, and it matches what is observed experimentally.

4.3 Step 2: identify the “local” regime (nearly diagonal kernel)

The kernel

$$K_{\mu_0}(x, x') = \int_{\mathbb{R}^d} \varphi_1(\langle \theta, x \rangle) \varphi_1(\langle \theta, x' \rangle) \mu_0(d\theta)$$

controls how strongly the update at x depends on data at X . If K_{μ_0} is localized (as in many wavelet constructions), then in a suitable grid setting one can be close to a diagonal-dominant regime,

$$K_{\mu_0}(x_i, x_j) \approx 0 \quad \text{for } i \neq j,$$

so each location x_i is driven primarily by its own residual and gating. In the ideal limit $K_{\mu_0}(x_i, x_j) = \delta_{ij}$, the dynamics decouple across i , and specialization can be studied pointwise in x . The general case can then be treated as a perturbation of this local regime.

4.4 Step 3: make reinforcement into a sign-coherent inequality

The essential nonlinear feedback comes from the ReLU gate inside B_k :

$$B_k(t; X) = \mathbb{E}_{C_2}[L_{C_2,k} w_2(t, C_2) \mathbb{1}\{H_2(t, C_2; X) > 0\}], \quad H_2(t, c_2; X) = \sum_{j=1}^r L_{c_2,j} f_j(t, X).$$

To prove a monotone reinforcement effect (“a larger f_j tends to increase its own future growth”), one typically restricts to a sign-coherent regime where:

- $L_{c_2,k} \geq 0$ on the relevant c_2 -mass (or has controlled sign changes),
- $w_2(t, c_2) \geq 0$ (or is aligned with $L_{c_2,\cdot}$),
- the residual signal $d_L(Z; W(t))$ has a stable sign on the region of X currently being fit (for instance locally in time or under one-sided targets).

Under such conditions, increasing $f_j(t, X)$ increases $H_2(t, C_2; X)$ on the set where $L_{C_2,j} > 0$, hence (weakly) increases the gate $\mathbb{1}\{H_2 > 0\}$, hence increases $B_j(t; X)$, closing a positive feedback loop through (??).

4.5 Step 4: prove specialization via relative-growth (replicator-style) dynamics

In a local regime (Step 2), define at fixed x the log-ratio

$$R_{k\ell}(t, x) \equiv \log \frac{|f_k(t, x)| + \varepsilon}{|f_\ell(t, x)| + \varepsilon}.$$

The objective is to show that for most (k, ℓ, x) one has a differential inequality of the form

$$\partial_t R_{k\ell}(t, x) \geq \text{Adv}_{k\ell}(t, x) - \text{Err}(t, x),$$

where the “advantage” term $\text{Adv}_{k\ell}$ is increasing in the current dominance of channel k at x due to the gating in B_k , and Err is controlled by off-diagonal kernel interactions and bad-event tails. Such inequalities yield exponential separation of ratios and thus winner-take-all channel shares $s_k(t, x)$.

4.6 Step 4.5: log-ratio calculus (exact identity) and a sufficient condition

Fix x and define $u_k(t, x) \equiv |f_k(t, x)| + \varepsilon$. Whenever $f_k(\cdot, x)$ is absolutely continuous in t , the map $u_k(\cdot, x)$ is also absolutely continuous and, for almost every t ,

$$\partial_t u_k(t, x) = \text{sign}(f_k(t, x)) \partial_t f_k(t, x) \quad \text{whenever } f_k(t, x) \neq 0,$$

and any choice in the interval $[-|\partial_t f_k(t, x)|, |\partial_t f_k(t, x)|]$ is a valid subderivative when $f_k(t, x) = 0$. In particular, for almost every t one has the exact log-ratio identity

$$\partial_t R_{k\ell}(t, x) = \frac{\text{sign}(f_k(t, x)) \partial_t f_k(t, x)}{|f_k(t, x)| + \varepsilon} - \frac{\text{sign}(f_\ell(t, x)) \partial_t f_\ell(t, x)}{|f_\ell(t, x)| + \varepsilon}. \quad (10)$$

Local decoupling and per-capita gains. In the ideal local regime $K_{\mu_0}(x, X) \approx K_{\mu_0}(x, x) \mathbb{1}\{X = x\}$ (or in the diagonal empirical-kernel approximation at a training point $x = x_i$), (??) reduces to

$$\partial_t f_k(t, x) \approx -\xi_1(t) K_{\mu_0}(x, x) \left(d_x(t) B_k(t; x) \right), \quad (11)$$

where $d_x(t)$ abbreviates the local residual signal (e.g. $d_L((x, y); W(t))$ in the one-point reduction). Plugging (??) into (??) motivates the *per-capita gain*

$$\gamma_k(t, x) \equiv \xi_1(t) K_{\mu_0}(x, x) \frac{-d_x(t) \text{sign}(f_k(t, x)) B_k(t; x)}{|f_k(t, x)| + \varepsilon},$$

so that $\partial_t R_{k\ell}(t, x) \approx \gamma_k(t, x) - \gamma_\ell(t, x)$. Thus, specialization at x is implied by a persistent gap in per-capita gains.

Lemma 6 (A sufficient per-capita gap condition). *Fix x and suppose there exist $\delta > 0$ and an index j such that on a time interval $[t_0, t_1]$,*

$$\gamma_j(t, x) \geq \gamma_k(t, x) + \delta \quad \text{for all } k \neq j \text{ and almost every } t \in [t_0, t_1].$$

Then for every $k \neq j$,

$$R_{jk}(t_1, x) \geq R_{jk}(t_0, x) + \delta (t_1 - t_0),$$

and hence the share of channel j at x increases relative to all other channels over $[t_0, t_1]$.

Example: the one-sparse spike regime. In the one-spike reduction (??)–(??) under the hypotheses of Lemma ??, we have $f(t) = a(t)e_j$ with $a(t)$ nondecreasing and $B_k(t) = 0$ for $k \neq j$. Then for $k \neq j$ the denominator $|f_k| + \varepsilon = \varepsilon$ is constant while $|f_j| + \varepsilon = a(t) + \varepsilon$ increases, so the log-ratio satisfies

$$\partial_t R_{jk}(t) = \frac{\partial_t a(t)}{a(t) + \varepsilon} = \frac{-\xi_1(t) d(t) B_j(t)}{a(t) + \varepsilon} > 0,$$

illustrating concretely how reinforcement yields monotone growth of log-ratios in a spike-dominated regime.

Remark 6 (The “minus sign” in $\partial_t f_k = -\xi_1 d B_k$ does not break specialization). *The minus sign is simply the gradient-descent direction. It does not turn reinforcement into anti-reinforcement by itself; what matters is the sign of the residual signal $d_x(t)$. For instance, for square loss $\mathcal{L}(y, \hat{y}) = \frac{1}{2}(\hat{y} - y)^2$, one has $d_x(t) = \hat{y}(x; W(t)) - y(x)$, which typically changes sign when the prediction crosses the target. Thus the spike amplitude $|f_j(t, x)|$ need not increase monotonically for all times. The robust specialization statement is instead about relative dominance: log-ratios $R_{jk}(t, x)$ can grow even during phases where all $|f_k(t, x)|$ decrease, because R_{jk} tracks per-capita gains (Step 4.5) rather than absolute growth.*

4.7 Step 4.6: effective sharpness along the spike direction (one-point square loss)

In a one-point reduction, one can make the intuition “the descent direction is clear” fully quantitative via a Lyapunov identity. Assume a single training pair $(x_\star, y_\star)$ with square loss $\mathcal{L} = \frac{1}{2}(\hat{y} - y_\star)^2$, and that we are on the one-sparse manifold of Lemma ?? so that $f(t) = a(t)e_j$ with $a(t) \geq 0$ and $\hat{y}(t) = a(t)B_j(t)$. Define the residual $d(t) \equiv \hat{y}(t) - y_\star$ and the scalar loss $\ell(t) \equiv \frac{1}{2}d(t)^2$.

Lemma 7 (Lyapunov identity and an “effective curvature” coefficient). *Assume the one-spike system (??)–(??), Assumption ??, and the one-sparse manifold $f(t) = a(t)e_j$ with $a(t) \geq 0$ on a time interval where the gate is $\mathbb{1}\{L_{C_2,j} > 0\}$. For square loss, the residual satisfies the closed scalar differential equation*

$$\partial_t d(t) = -d(t) \left(\xi_1(t) B_j(t)^2 + \xi_2(t) a(t)^2 S_j \right), \quad S_j \equiv \mathbb{E}[L_{C_2,j}^2 \mathbb{1}\{L_{C_2,j} > 0\}]. \quad (12)$$

Consequently the loss decreases according to

$$\partial_t \ell(t) = -2 \ell(t) \left(\xi_1(t) B_j(t)^2 + \xi_2(t) a(t)^2 S_j \right) \leq 0. \quad (13)$$

Proof deferred to appendix.

Remark 7 (What “sharpness” means here). *The coefficient $\lambda(t) \equiv \xi_1(t) B_j(t)^2 + \xi_2(t) a(t)^2 S_j$ plays the role of an effective curvature along the active spike direction. If $\lambda(t) \geq \lambda_0 > 0$ on an interval, then (??) gives exponential decay $\ell(t) \leq \ell(t_0) e^{-2\lambda_0(t-t_0)}$. This is a rigorous version of “the landscape is sharp in the useful direction”. It should not be confused with global strong convexity: low-rank / bilinear parametrizations can retain flat directions (e.g. scaling symmetries) even while being very sharp along the spike-aligned direction.*

4.8 Step 5: connect specialization to spikes using the feature dictionary

Specialization alone does not force spikes; the spike shape is controlled by (μ_0, φ_1) through K_{μ_0} . For wavelet-like features (or other localized dictionaries), channel specialization on disjoint regions implies that different channels are effectively solving different local regression problems, which can produce localized high-amplitude fits (spikes) at different locations. Thus a practical strategy is:

- prove channel specialization (Step 4) in a sign-coherent, local-kernel regime,
- then quantify how localization properties of K_{μ_0} convert specialization into spike-like structure.

4.9 Step 5.1: closed-form dynamics in the kernel (NTK) regime and frequency-by-frequency learning

This subsection records an explicit “solved” dynamics in the regime where the learning dynamics reduce to a *linear* kernel gradient flow with a *fixed* Gram operator. This is the cleanest mathematical lens for understanding learning of localized targets (two-sided steps) or pure frequencies (cosines).

Empirical (finite-sample) kernel dynamics. Fix training points $x_1, \dots, x_n$ with targets $y = (y_1, \dots, y_n)^\top$ and let $\hat{y}(t) \in \mathbb{R}^n$ be the vector of current predictions $\hat{y}_i(t) \equiv \hat{y}(x_i; W(t))$. In a standard NTK/random-features linearization (kernel frozen at initialization), the squared-loss gradient flow closes as

$$\partial_t \hat{y}(t) = -\eta(t) K (\hat{y}(t) - y), \quad (14)$$

where $K \in \mathbb{R}^{n \times n}$ is a fixed positive semidefinite Gram matrix (typically $K_{ij} = K_{\mu_0}(x_i, x_j)$ up to a scaling depending on parametrization) and $\eta(t) \geq 0$ is an effective learning-rate factor.

Lemma 8 (Matrix exponential solution). *Assume K does not depend on t and η is integrable on $[0, \infty)$. Define the effective time $\tau(t) \equiv \int_0^t \eta(s) ds$. Then the solution of (??) is*

$$\hat{y}(t) - y = \exp(-\tau(t) K) (\hat{y}(0) - y). \quad (15)$$

Equivalently, if $K = U \text{diag}(\lambda_1, \dots, \lambda_n) U^\top$ with $\lambda_i \geq 0$, then the mode coefficients $\alpha_i(t) \equiv u_i^\top (\hat{y}(t) - y)$ satisfy

$$\alpha_i(t) = e^{-\lambda_i \tau(t)} \alpha_i(0), \quad i = 1, \dots, n. \quad (16)$$

Proof deferred to appendix.

Continuous kernel flow and Fourier modes (translation-invariant case). To connect directly to “frequency learning”, suppose $x \in \mathbb{R}$ (or $x \in \mathbb{T}$) and the data distribution is uniform so that the kernel operator is approximately translation-invariant: $K_{\mu_0}(x, x') = \kappa(x - x')$ for a function κ . Then the squared-loss kernel gradient flow takes the form

$$\partial_t \hat{y}(t, x) = -\eta(t) \int \kappa(x - x') (\hat{y}(t, x') - y(x')) dx'. \quad (17)$$

Taking Fourier transforms gives a fully decoupled dynamics: for each frequency ω ,

$$\partial_t \widehat{(\hat{y} - y)}(t, \omega) = -\eta(t) \widehat{\kappa}(\omega) \widehat{(\hat{y} - y)}(t, \omega), \quad \widehat{(\hat{y} - y)}(t, \omega) = e^{-\widehat{\kappa}(\omega) \tau(t)} \widehat{(\hat{y} - y)}(0, \omega). \quad (18)$$

Thus each frequency decays exponentially at a rate set by the kernel spectrum $\widehat{\kappa}(\omega)$.

Cosine target (single frequency). If $y(x) = A \cos(\omega_0 x)$ and $\hat{y}(0, \cdot)$ has no ω_0 component, then (??) implies

$$\hat{y}(t, x) = A(1 - e^{-\widehat{\kappa}(\omega_0) \tau(t)}) \cos(\omega_0 x),$$

so the learning time-scale of frequency ω_0 is exactly $[\widehat{\kappa}(\omega_0)]^{-1}$ (up to the effective time τ). Frequency bias is therefore determined by monotonicity of $\widehat{\kappa}(\omega)$ in $|\omega|$ (many smooth kernels have $\widehat{\kappa}(\omega)$ decreasing in $|\omega|$, yielding slow learning of high frequencies).

Two-sided step target (localized jump). Consider a two-sided step on $\mathbb{T}$ or on a bounded interval, e.g.

$$y(x) = A(\mathbb{1}\{|x - x_0| \leq \delta\} - \mathbb{1}\{|x - x_1| \leq \delta\}),$$

which is the difference of two localized bumps (“two-sided” in the sense of two opposite signed supports). Its Fourier coefficients satisfy $|\widehat{y}(\omega)| \asymp |\omega|^{-1}$ for large $|\omega|$ (jump discontinuities generate a $1/\omega$ tail). By (??), each mode decays as $e^{-\widehat{\kappa}(\omega) \tau(t)}$. Hence the time needed to resolve sharp edges is controlled by how quickly $\widehat{\kappa}(\omega)$ decays with $|\omega|$:

- if $\widehat{\kappa}(\omega)$ decays rapidly (very smooth kernels), high frequencies are learned slowly and the step is learned by progressive sharpening;
- if $\widehat{\kappa}(\omega)$ is flatter across a wide band (more localized dictionaries, less smooth kernels), the step “locks in” quickly, producing spike-like localization.

Example: 4-point “two-sided step” dataset $\{0, \frac{1}{3}, \frac{2}{3}, 1\}$ with labels $(0, A, -A, 0)$. Let $x_1 = 0, x_2 = \frac{1}{3}, x_3 = \frac{2}{3}, x_4 = 1$ and $y = (0, A, -A, 0)^\top$. Assume a symmetric distance kernel $K_{ij} = \kappa(|x_i - x_j|)$ (this includes the NNGP kernel in many isotropic setups restricted to a 1D grid). Set

$$k_0 \equiv \kappa(0), \quad k_1 \equiv \kappa\left(\frac{1}{3}\right), \quad k_2 \equiv \kappa\left(\frac{2}{3}\right), \quad k_3 \equiv \kappa(1),$$

so the Gram matrix is the 4×4 Toeplitz matrix

$$K = \begin{pmatrix} k_0 & k_1 & k_2 & k_3 \\ k_1 & k_0 & k_1 & k_2 \\ k_2 & k_1 & k_0 & k_1 \\ k_3 & k_2 & k_1 & k_0 \end{pmatrix}.$$

Consider the kernel flow (??) with $\hat{y}(0) = 0$, hence $e(0) \equiv \hat{y}(0) - y = -y$. Because K is Toeplitz and y is antisymmetric, the dynamics stay in the 2D antisymmetric subspace

$$e(t) = (a(t), b(t), -b(t), -a(t))^\top, \quad \text{equivalently} \quad \hat{y}(t) = (a(t), A + b(t), -A - b(t), -a(t))^\top.$$

In this subspace the error dynamics close as a 2×2 linear system (in effective time $\tau(t) = \int_0^t \eta(s) ds$):

$$\partial_\tau \begin{pmatrix} a \\ b \end{pmatrix} = -M \begin{pmatrix} a \\ b \end{pmatrix}, \quad M = \begin{pmatrix} k_0 - k_3 & k_1 - k_2 \\ k_1 - k_2 & k_0 - k_1 \end{pmatrix}, \quad (a(0), b(0)) = (0, -A). \quad (19)$$

Write

$$m \equiv \frac{\text{tr}(M)}{2} = k_0 - \frac{k_1 + k_3}{2}, \quad d \equiv \frac{M_{11} - M_{22}}{2} = \frac{k_1 - k_3}{2}, \quad o \equiv M_{12} = k_1 - k_2, \quad s \equiv \sqrt{d^2 + o^2}.$$

Then the solution is explicit:

$$a(t) = A e^{-m\tau(t)} \frac{o}{s} \sinh(s\tau(t)), \quad b(t) = -A e^{-m\tau(t)} \left(\cosh(s\tau(t)) + \frac{d}{s} \sinh(s\tau(t)) \right). \quad (20)$$

In particular, the two nonzero labels are learned through the single scalar $\hat{y}(t, x_2) = A + b(t)$ (and $\hat{y}(t, x_3) = -A - b(t)$), while the endpoints take the coupling-induced values $\hat{y}(t, x_1) = a(t)$ and $\hat{y}(t, x_4) = -a(t)$.

Computing (k_0, k_1, k_2, k_3) explicitly from a closed-form NNGP kernel (ReLU). To obtain concrete numerical rates, one can plug in a closed-form NNGP kernel. Assume:

- $\theta \sim \mathcal{N}(0, I_d)$ and $x, x' \in \mathbb{R}^d$ are embedded/normalized so that $\|x\| = \|x'\| = 1$ and $\rho \equiv \langle x, x' \rangle \in [-1, 1]$,
- $\varphi_1(u) = u_+$ (ReLU first-layer features).

Then the NNGP (feature) kernel has the well-known closed form

$$K_{\mu_0}(x, x') = \kappa_{\text{ReLU}}(\rho) \equiv \frac{1}{2\pi} \left(\sqrt{1 - \rho^2} + (\pi - \arccos \rho) \rho \right), \quad \rho = \langle x, x' \rangle. \quad (21)$$

If the 1D coordinate $t \in [0, 1]$ is embedded on the unit circle, $x(t) = (\cos(2\pi t), \sin(2\pi t))$, then $\rho(t, t') = \cos(2\pi(t - t'))$ depends only on the distance on the circle. For the grid $\{0, \frac{1}{3}, \frac{2}{3}, 1\}$ one has

$$\rho(0, 1) = 1, \quad \rho(0, \frac{1}{3}) = \rho(0, \frac{2}{3}) = \cos(\frac{2\pi}{3}) = -\frac{1}{2},$$

so $k_0 = k_3 = \kappa_{\text{ReLU}}(1) = \frac{1}{2}$ and $k_1 = k_2 = \kappa_{\text{ReLU}}(-\frac{1}{2})$ with

$$k_1 = k_2 = \frac{1}{2\pi} \left(\frac{\sqrt{3}}{2} - \frac{\pi}{6} \right) = \frac{\sqrt{3}}{4\pi} - \frac{1}{12}. \quad (22)$$

In particular $k_1 = k_2$ and $k_0 = k_3$, hence the 2×2 system (??) reduces to the scalar ODE

$$\partial_\tau b(\tau) = -(k_0 - k_1) b(\tau), \quad b(0) = -A,$$

so

$$\hat{y}(t, x_2) = A + b(t) = A \left(1 - e^{-(k_0 - k_1)\tau(t)}\right), \quad \hat{y}(t, x_3) = -A - b(t) = -A \left(1 - e^{-(k_0 - k_1)\tau(t)}\right), \quad (23)$$

with the explicit rate

$$k_0 - k_1 = \frac{1}{2} - \left(\frac{\sqrt{3}}{4\pi} - \frac{1}{12}\right) = \frac{7}{12} - \frac{\sqrt{3}}{4\pi}.$$

Lemma 9 (Exact learning curve for the two-sided step on the 4-point grid in the ReLU NNGP kernel regime). *In the 4-point dataset $x_1 = 0$, $x_2 = \frac{1}{3}$, $x_3 = \frac{2}{3}$, $x_4 = 1$ with $y = (0, A, -A, 0)^\top$, assume the fixed-kernel gradient flow (??) with $\hat{y}(0) = 0$ and the ReLU NNGP kernel (??) under the unit-circle embedding $x(t) = (\cos(2\pi t), \sin(2\pi t))$. Then the predictions at the two nonzero-label points satisfy the explicit learning curves*

$$\hat{y}(t, x_2) = A \left(1 - e^{-\lambda \tau(t)}\right), \quad \hat{y}(t, x_3) = -A \left(1 - e^{-\lambda \tau(t)}\right), \quad \lambda = \frac{7}{12} - \frac{\sqrt{3}}{4\pi},$$

and $\hat{y}(t, x_1) = \hat{y}(t, x_4) = 0$ for all t . In particular, the (normalized) step amplitude

$$g(t) \equiv \frac{\hat{y}(t, x_2)}{A} = 1 - e^{-\lambda \tau(t)}$$

increases monotonically from 0 to 1 at an exponential rate controlled by λ .

Remark 8 (How to view this as a “curve of f_k ” in the spike dictionary picture). *In Step 5.2, in a diagonal-dominant regime the channel functions are superpositions of kernel bumps. On the 4-point step dataset, the target is supported on the two interior points $\{x_2, x_3\}$. If one is in a regime where the two channels specialize so that channel 1 carries the positive spike and channel 2 carries the negative spike, a natural idealized representation is*

$$f_1(t, x) \approx A g(t) \frac{K_{\mu_0}(x, x_2)}{K_{\mu_0}(x_2, x_2)}, \quad f_2(t, x) \approx -A g(t) \frac{K_{\mu_0}(x, x_3)}{K_{\mu_0}(x_3, x_3)},$$

with the same scalar amplitude curve $g(t) = 1 - e^{-\lambda \tau(t)}$. For the ReLU NNGP kernel under the unit-circle embedding, this becomes an explicit bump shape on $t \in [0, 1]$: for $x = x(t)$ and $x_p = x(t_p)$,

$$K_{\mu_0}(x(t), x(t_p)) = \kappa_{\text{ReLU}}(\cos(2\pi(t - t_p))),$$

so the entire shape of $f_1(t, \cdot)$ and $f_2(t, \cdot)$ is given in closed form by κ_{ReLU} and the scalar amplitude $g(t)$. The remaining mathematical task (handled by Step 4.5–Step 5.3) is to justify this specialization from the nonlinear mean-field dynamics, rather than postulating it.

Remark 9 (High-dimensional “almost diagonal” limit gives an even simpler closed form). *If off-diagonal kernel entries are approximately equal (a common high-dimensional concentration pattern after centering), $k_1 \approx k_2 \approx k_3 \approx k_{\text{off}}$, then $o \approx 0$ and $d \approx 0$, hence $a(t) \approx 0$ and*

$$\hat{y}(t, x_2) \approx A \left(1 - e^{-(k_0 - k_{\text{off}})\tau(t)}\right), \quad \hat{y}(t, x_3) \approx -A \left(1 - e^{-(k_0 - k_{\text{off}})\tau(t)}\right),$$

so the “frequency” mode $(0, 1, -1, 0)$ is learned at rate $k_0 - k_{\text{off}}$. If moreover the features are mean-zero / centered so that $k_{\text{off}} \approx 0$, then the four points essentially decouple and the learning rate is $\approx k_0$.

Remark 10 (What is “fully solved” and what remains model-specific). *The formulas (??)–(??) are an exact closed-form solution once one commits to a fixed kernel (NTK regime). The model-specific work is to justify that, in the MMNN low-rank mean-field dynamics, the relevant Gram operator is close to a fixed K_{μ_0} (or to quantify how gating makes it time-dependent). In the diagonal-dominant regime $K \approx \text{diag}(K_{11}, \dots, K_{nn})$, (??) reduces to n independent scalar ODEs and step-like targets behave as collections of nearly independent “spikes”.*

4.10 Step 5.2: solving the mean-field w_1 -PDE explicitly for finitely-supported data (and reducing spike learning to scalar ODEs)

We now go back from the fixed-kernel linearization (Step 5.1) to the actual mean-field equations for (w_1, w_2) . The key point is that the w_1 evolution is *linear* in feature space once $B_k(t; \cdot)$ and the residual signal are fixed, so for finitely-supported data one can solve $\bar{w}_{1,k}$ and hence f_k *exactly* as a superposition of kernel bumps.

Finite-support data and scalar coefficients. Assume the input marginal is supported on finitely many points $x^{(1)}, \dots, x^{(m)}$:

$$X \in \{x^{(1)}, \dots, x^{(m)}\}, \quad \mathbb{P}(X = x^{(p)}) = \pi_p, \quad p = 1, \dots, m,$$

with corresponding targets $y^{(p)}$. For each location $x^{(p)}$, write

$$d_p(t) \equiv d_L((x^{(p)}, y^{(p)}); W(t)), \quad B_{k,p}(t) \equiv B_k(t; x^{(p)}).$$

Then (??) becomes the explicit ODE (pointwise in θ)

$$\partial_t \bar{w}_{1,k}(t, \theta) = -\xi_1(t) \sum_{p=1}^m \pi_p d_p(t) B_{k,p}(t) \varphi_1(\langle \theta, x^{(p)} \rangle), \quad k = 1, \dots, r. \quad (24)$$

Integrating (??) in time yields the closed form

$$\bar{w}_{1,k}(t, \theta) = \bar{w}_{1,k}(0, \theta) - \sum_{p=1}^m \varphi_1(\langle \theta, x^{(p)} \rangle) \Gamma_{k,p}(t), \quad \Gamma_{k,p}(t) \equiv \int_0^t \xi_1(s) \pi_p d_p(s) B_{k,p}(s) ds. \quad (25)$$

Plugging (??) into the definition of f_k gives an explicit superposition in input space:

$$f_k(t, x) = f_k(0, x) - \sum_{p=1}^m \Gamma_{k,p}(t) K_{\mu_0}(x, x^{(p)}), \quad k = 1, \dots, r. \quad (26)$$

Thus, in finitely-supported settings, the *spike shape* of each channel is completely controlled by the dictionary kernel K_{μ_0} , while all learning dynamics reduce to the scalar coefficients $\Gamma_{k,p}(t)$.

One-spike reduction (exact spike shape). If $m = 1$ with $x^{(1)} = x_\star$ (a single spike location), then (??) becomes

$$f_k(t, x) = f_k(0, x) - \Gamma_{k,1}(t) K_{\mu_0}(x, x_\star).$$

In particular, if $f_k(0, \cdot) \equiv 0$ then $f_k(t, \cdot)$ is *exactly* a kernel bump at $x_\star$ for all t , and “spike learning” reduces to solving $\Gamma_{k,1}(t)$ (equivalently $f_k(t, x_\star)$). This makes precise why Step 5 depends on the localization of K_{μ_0} : the dynamics can only create spikes to the extent that $K_{\mu_0}(x, x_\star)$ is itself localized.

Two-sided step as a two-spike reduction. A two-sided step can be idealized as two opposite-signed locations, $m = 2$, e.g. $y^{(1)} = +A$ at $x^{(1)} = x_0$ and $y^{(2)} = -A$ at $x^{(2)} = x_1$. Then each channel has the exact representation

$$f_k(t, x) = f_k(0, x) - \Gamma_{k,1}(t) K_{\mu_0}(x, x_0) - \Gamma_{k,2}(t) K_{\mu_0}(x, x_1).$$

If K_{μ_0} is diagonal-dominant on $\{x_0, x_1\}$ (i.e. $|K_{\mu_0}(x_0, x_1)| \ll K_{\mu_0}(x_0, x_0)$), the two coefficients $(\Gamma_{k,1}, \Gamma_{k,2})$ are only weakly coupled through the residuals, and one obtains two almost-independent “local” spike dynamics (one near x_0 and one near x_1).

How to read “instability at initialization” in this formulation. In this finite-support reduction, “initialization is unstable” toward specialization at a location $x^{(p)}$ means: starting from small but unequal coefficients $\Gamma_{k,p}(0)$ across k , the log-ratios

$$R_{k\ell}(t, x^{(p)}) = \log \frac{|f_k(t, x^{(p)})| + \varepsilon}{|f_\ell(t, x^{(p)})| + \varepsilon}$$

grow because the per-capita gains (Step 4.5) generate a persistent gap. In the exact one-sparse spike regime of Lemma ??, this instability is extreme: all non-dominant channels satisfy $B_k(t) = 0$ and hence $\Gamma_{k,p}(t) \equiv 0$, while the dominant channel has $B_j(t) > 0$ and $\Gamma_{j,p}(t)$ grows in the descent direction as long as the residual has the appropriate sign. The remaining (hard) step beyond this exact regime is a stability estimate: if one channel is only *approximately* dominant, then B_k (hence $\Gamma_{k,p}$) is small but nonzero for $k \neq j$, yielding quantitative log-ratio growth rather than exact decoupling.

Case $r = 2$, two-spike setting, and “slight dominance” at one spike. Take $r = 2$ channels and $m = 2$ spike locations $x^{(1)} = x_0$ and $x^{(2)} = x_1$ with opposite targets (two-sided step). Assume a diagonal-dominant local regime so that $K_{\mu_0}(x_0, x_1)$ is negligible compared to $K_{\mu_0}(x_0, x_0)$ and $K_{\mu_0}(x_1, x_1)$. If moreover $f_k(0, \cdot) \equiv 0$, then evaluating (??) at x_0 gives the local approximation

$$f_k(t, x_0) \approx -\Gamma_{k,1}(t) K_{\mu_0}(x_0, x_0), \quad k = 1, 2, \quad (27)$$

so “channel 1 is slightly dominant at the spike x_0 ” is the quantitative condition

$$|f_1(t, x_0)| > |f_2(t, x_0)| \quad \Longleftrightarrow \quad |\Gamma_{1,1}(t)| > |\Gamma_{2,1}(t)|$$

in the local approximation (??). To conclude that this dominance is self-reinforcing, one needs a *stability* form of the one-sparse mechanism: the backprop signals satisfy a bound of the form $|B_{2,1}(t)| \ll |B_{1,1}(t)|$ whenever $|f_2(t, x_0)| \ll |f_1(t, x_0)|$.

Lemma 10 (A local log-ratio growth criterion for $r = 2$). *Fix the spike location x_0 and assume the local reduction (??) holds there:*

$$\partial_t f_k(t, x_0) = -\xi_1(t) K_{\mu_0}(x_0, x_0) d_0(t) B_{k,1}(t), \quad k = 1, 2,$$

where $d_0(t)$ is the residual signal at x_0 and $B_{k,1}(t) = B_k(t; x_0)$. Assume there exists a constant $\rho \in [0, 1)$ such that on a time interval I ,

$$|B_{2,1}(t)| \leq \rho \frac{|f_2(t, x_0)| + \varepsilon}{|f_1(t, x_0)| + \varepsilon} |B_{1,1}(t)| \quad \text{for all } t \in I. \quad (28)$$

If moreover $-d_0(t) \geq 0$ and $B_{1,1}(t)$ has the same sign as $f_1(t, x_0)$ on I , then the log-ratio

$$R_{12}(t, x_0) = \log \frac{|f_1(t, x_0)| + \varepsilon}{|f_2(t, x_0)| + \varepsilon}$$

satisfies $\partial_t R_{12}(t, x_0) \geq (1 - \rho) \xi_1(t) K_{\mu_0}(x_0, x_0) (-d_0(t)) \frac{|B_{1,1}(t)|}{|f_1(t, x_0)| + \varepsilon} \geq 0$ for a.e. $t \in I$. In particular, any strict dominance at $t = t_0 \in I$ cannot be lost on I and is amplified whenever $|B_{1,1}|$ is not too small.

Proof deferred to appendix.

Lemma 11 (Two-sided step (two spikes): log-ratio growth at each spike under local decoupling). Assume the two-sided step setting of Step 5.2 with $m = 2$ spike locations $x^{(1)} = x_0$ and $x^{(2)} = x_1$, and take $r = 2$ channels. Assume a diagonal-dominant regime so that the local reductions hold at each spike: for $p \in \{0, 1\}$ and $k \in \{1, 2\}$,

$$\partial_t f_k(t, x_p) = -\xi_1(t) K_{\mu_0}(x_p, x_p) d_p(t) B_{k,p+1}(t),$$

where $d_p(t)$ is the residual signal at x_p and $B_{k,p+1}(t) = B_k(t; x^{(p+1)})$. Suppose that on a time interval I :

- $-d_0(t) \geq 0$ and $B_{1,1}(t)$ has the same sign as $f_1(t, x_0)$, and there exists $\rho_0 \in [0, 1)$ such that

$$|B_{2,1}(t)| \leq \rho_0 \frac{|f_2(t, x_0)| + \varepsilon}{|f_1(t, x_0)| + \varepsilon} |B_{1,1}(t)|,$$

- $-d_1(t) \geq 0$ and $B_{2,2}(t)$ has the same sign as $f_2(t, x_1)$, and there exists $\rho_1 \in [0, 1)$ such that

$$|B_{1,2}(t)| \leq \rho_1 \frac{|f_1(t, x_1)| + \varepsilon}{|f_2(t, x_1)| + \varepsilon} |B_{2,2}(t)|.$$

Then the log-ratios at the two spikes satisfy, for a.e. $t \in I$,

$$\partial_t R_{12}(t, x_0) \geq (1 - \rho_0) \xi_1(t) K_{\mu_0}(x_0, x_0) (-d_0(t)) \frac{|B_{1,1}(t)|}{|f_1(t, x_0)| + \varepsilon} \geq 0,$$

and

$$\partial_t R_{21}(t, x_1) \geq (1 - \rho_1) \xi_1(t) K_{\mu_0}(x_1, x_1) (-d_1(t)) \frac{|B_{2,2}(t)|}{|f_2(t, x_1)| + \varepsilon} \geq 0,$$

where $R_{21}(t, x_1) = \log \frac{|f_2(t, x_1)| + \varepsilon}{|f_1(t, x_1)| + \varepsilon}$. In particular, any initial strict dominance of channel 1 at x_0 and of channel 2 at x_1 cannot be lost on I and is amplified whenever the corresponding $|B|$ terms are not too small.

Remark 11 (Why this is the “step setup” version of Lemma ??). Lemma ?? gives a one-spike sufficient condition for winner-take-all (ratio growth) at a single location. Lemma ?? is the direct two-spike specialization: under local decoupling, one applies the same log-ratio calculus at both spike locations. The remaining model-specific work is therefore reduced to verifying the stability-type inequalities on B (the “ $\rho < 1$ ” conditions) from the gating dynamics and the mixing structure.

Remark 12 (How ?? can arise from a half-space gate). Condition ?? is an abstract “approximate one-sparse stability” inequality. In the random-sign mixing picture, it is expected to hold when the gate $\mathbb{1}\{\langle L_{C_2}, f(t, x_0) \rangle > 0\}$ is largely determined by the dominant coordinate (here channel 1), so that the effective selection of c_2 -indices creates a strong conditional bias for $B_{1,1}$ but only a weak (order $|f_2|/|f_1|$) bias for $B_{2,1}$. Making this fully rigorous requires quantitative control of the dependence of $w_2(t, C_2)$ on the small coordinate through ??; in the exact one-sparse regime of Lemma ?? one has $\rho = 0$.

4.11 Step 5.3: the full local $r = 2$ spike dynamics for (f_1, f_2) , per-capita gains, and what can (and cannot) be “solved”

This subsection writes the complete closed system at a spike location in the local regime and clarifies which structural assumptions allow an explicit solution or yield (non-)specialization.

Local two-channel dynamics at one spike location. Fix a spike location x_0 and assume a diagonal-dominant local regime so that $K_{\mu_0}(x_0, X) \approx K_{\mu_0}(x_0, x_0) \mathbb{1}\{X = x_0\}$. Write

$$u_1(t) \equiv f_1(t, x_0), \quad u_2(t) \equiv f_2(t, x_0), \quad u(t) = (u_1(t), u_2(t)) \in \mathbb{R}^2,$$

and let $L_{c_2} = (L_{c_2,1}, L_{c_2,2}) \in \mathbb{R}^2$ be the frozen mixing vector for index c_2 . Define the preactivation and gate at x_0 :

$$S(t, c_2) \equiv \langle L_{c_2}, u(t) \rangle, \quad G(t, c_2) \equiv \mathbb{1}\{S(t, c_2) > 0\}.$$

For ReLU top activation $\varphi_2(u) = u_+$ and square loss at (x_0, y_0) , the local residual is $d_0(t) \equiv \hat{y}(t, x_0) - y_0$ with

$$\hat{y}(t, x_0) = \mathbb{E}_{C_2}[w_2(t, C_2) S(t, C_2) G(t, C_2)]. \quad (29)$$

The backprop signals are

$$B_k(t; x_0) = \mathbb{E}_{C_2}[L_{C_2,k} w_2(t, C_2) G(t, C_2)], \quad k = 1, 2. \quad (30)$$

Then the (local) channel dynamics are

$$\partial_t u_k(t) = -\xi_1(t) K_{\mu_0}(x_0, x_0) d_0(t) B_k(t; x_0), \quad k = 1, 2, \quad (31)$$

and the top-layer mean-field dynamics are (finite-support reduction with one point)

$$\partial_t w_2(t, c_2) = -\xi_2(t) d_0(t) S(t, c_2) G(t, c_2). \quad (32)$$

Equations (??)–(??) are the complete closed ODE system for $(u(t), w_2(t, \cdot))$ at one spike location.

Per-capita gains and log-ratio dynamics. Define the log-ratio at x_0 ,

$$R_{12}(t, x_0) = \log \frac{|u_1(t)| + \varepsilon}{|u_2(t)| + \varepsilon}.$$

Combining (??) with (??) yields the exact identity (a.e. in t)

$$\partial_t R_{12}(t, x_0) = \xi_1(t) K_{\mu_0}(x_0, x_0) (-d_0(t)) \left(\frac{\text{sign}(u_1(t)) B_1(t; x_0)}{|u_1(t)| + \varepsilon} - \frac{\text{sign}(u_2(t)) B_2(t; x_0)}{|u_2(t)| + \varepsilon} \right). \quad (33)$$

Thus the whole winner-take-all question at a spike reduces to comparing the *per-capita gains*

$$\gamma_k(t, x_0) \equiv \xi_1(t) K_{\mu_0}(x_0, x_0) (-d_0(t)) \frac{\text{sign}(u_k(t)) B_k(t; x_0)}{|u_k(t)| + \varepsilon}, \quad k = 1, 2.$$

What is explicitly solvable.

- *Exact one-sparse manifold.* If $u_2(0) = 0$ and the mixing distribution is such that $B_2(t; x_0) = 0$ whenever $u_2(t) = 0$ (as in Lemma ??), then $u_2(t) \equiv 0$ and u_1 reduces to the scalar dynamics already solved in Lemma ?? and Lemma ??.
- *Fixed-kernel linearization.* If one replaces (B_1, B_2) by a linear function of u frozen at initialization, one recovers the matrix exponential solution of Step 5.1.

A useful warning: isotropic Gaussian mixing forces $B(t; x_0)$ to be parallel to $u(t)$ (no ratio dynamics). The half-space mechanism alone does *not* automatically create specialization when the mixing distribution is isotropic. The following lemma makes this precise.

Lemma 12 (Isotropic Gaussian mixing implies B is colinear with u). *Assume $L_{C_2} \sim \mathcal{N}(0, I_2)$ and that $w_2(t, C_2)$ is a measurable function of $S(t, C_2) = \langle L_{C_2}, u(t) \rangle$ only, i.e. $w_2(t, C_2) = g_t(S(t, C_2))$ for some function $g_t : \mathbb{R} \rightarrow \mathbb{R}$. Then for any $u(t) \neq 0$,*

$$(B_1(t; x_0), B_2(t; x_0)) = c(t) u(t)$$

for a scalar coefficient $c(t)$ depending on $u(t)$ only through $\|u(t)\|$ and on g_t through $\mathbb{E}[S g_t(S) \mathbb{1}\{S > 0\}]$. Consequently, if $u_1(t), u_2(t)$ have the same sign, then the bracket in (??) vanishes and $R_{12}(t, x_0)$ is constant (no winner-take-all).

Proof deferred to appendix.

Remark 13 (What this means for your experiments). *Lemma ?? explains why proving specialization requires more than “half-space gating + symmetry”: one needs either (i) non-isotropic / structured mixing L (e.g. dictionary-like rows), or (ii) multi-spike interactions (different gates at different x reshape w_2 in a way that breaks the $w_2 = g(S)$ symmetry at a single location), or (iii) higher-layer structure beyond the one-spike reduction. In such settings, the stability hypothesis (??) becomes plausible and yields strict log-ratio growth via Lemma ??.*

4.12 Step 5.4: seeding strict dominance by the first SGD step (channel-resolved NTK/linearization)

The previous steps show that *if* one has strict dominance at each spike location and a stability bound of the form (??), then winner-take-all amplification follows. This subsection explains a complementary (and often simpler) mechanism to obtain the *initial* strict dominance: the very first stochastic gradient step (finite-width, random initialization) typically produces a small but *random* channel imbalance at each spike. One then bootstraps: first-step imbalance $\Rightarrow$ log-ratio becomes positive, and afterwards the continuous-time arguments apply.

Discrete-time first step in the two-spike reduction. Consider the two-spike (two-sided step) dataset with $x^{(1)} = x_0$, $x^{(2)} = x_1$ and $y^{(1)} = +A$, $y^{(2)} = -A$, and take $r = 2$ channels. Assume a local/diagonal-dominant regime so that the evolution at each spike is governed primarily by the diagonal kernel values $K_{\mu_0}(x_p, x_p)$. Assume also that the initial prediction is close to 0 at both spikes, so the initial residuals satisfy

$$d_0(0) \approx -A, \quad d_1(0) \approx +A,$$

for square loss (underfitting at x_0 and overfitting at x_1 relative to 0).

Let $\eta > 0$ be a (small) SGD step size and consider the first-step update of the channel values at the spikes under a channel-resolved linearization (freeze the kernel K_{μ_0} and the backprop signals $B_{k,p}(0)$ at initialization):

$$f_k^+(x_p) \equiv f_k(0, x_p) - \eta \xi_1(0) \pi_p d_p(0) B_{k,p}(0) K_{\mu_0}(x_p, x_p), \quad k = 1, 2, \quad p \in \{0, 1\}, \quad (34)$$

where $\pi_p = \mathbb{P}(X = x_p)$ and $B_{k,p}(0) = B_k(0; x^{(p+1)})$. The precise update rule depends on the optimizer/minibatch, but (??) captures the leading-order effect: the first step creates channel amplitudes proportional to the random initialization of $B_{k,p}(0)$.

Lemma 13 (First-step strict dominance occurs with constant probability under exchangeable initialization). *Assume that at initialization the random variables $(|B_{1,1}(0)|, |B_{2,1}(0)|)$ are i.i.d. with a continuous distribution, and likewise $(|B_{1,2}(0)|, |B_{2,2}(0)|)$ are i.i.d., independent of the first pair. Assume $f_k(0, x_p) = 0$ and that (??) holds with $K_{\mu_0}(x_p, x_p) > 0$ and $\pi_p > 0$. Then:*

- at x_0 , $\mathbb{P}(|f_1^+(x_0)| > |f_2^+(x_0)|) = \frac{1}{2}$;
- at x_1 , $\mathbb{P}(|f_2^+(x_1)| > |f_1^+(x_1)|) = \frac{1}{2}$;
- consequently, the “correct assignment” event

$$\mathcal{E} \equiv \{|f_1^+(x_0)| > |f_2^+(x_0)|\} \cap \{|f_2^+(x_1)| > |f_1^+(x_1)|\}$$

has probability $\mathbb{P}(\mathcal{E}) = \frac{1}{4}$.

Moreover, on $\mathcal{E}$ the initial log-ratios are strictly positive:

$$R_{12}(0^+, x_0) = \log \frac{|f_1^+(x_0)| + \varepsilon}{|f_2^+(x_0)| + \varepsilon} > 0, \quad R_{21}(0^+, x_1) = \log \frac{|f_2^+(x_1)| + \varepsilon}{|f_1^+(x_1)| + \varepsilon} > 0.$$

Proof deferred to appendix.

Remark 14 (How this plugs into the continuous-time winner-take-all argument). *Lemma ?? is only a seeding result: it produces an initial strict dominance with constant probability from exchangeable random initialization. To conclude that the dominance is amplified over time, one then proves a stability bound of the form (??) (or its two-spike version in Lemma ??), which yields monotone log-ratio growth. This is exactly the bootstrap structure: first-step symmetry breaking $\Rightarrow$ ratios become positive $\Rightarrow$ stability implies ratios keep increasing.*

4.13 Step 6: what is realistically provable first

A plausible first theorem is a *symmetry-breaking / specialization* result: under sign-coherence and a diagonal-dominant kernel assumption, the symmetric state where channels are comparable becomes unstable and the dynamics drive $s_k(t, x)$ toward $\{0, 1\}$ (one dominant channel) for most x . This already captures the observed “different channels learn different slots” effect, and it is robust to small perturbations (off-diagonal kernel terms) via Grönwall and stability estimates.

5 Layer-wise spike sharpening and frequency growth across depth

(this section describes a proposed mechanism and proof targets; the key assumptions are not yet proved).

This section targets the statement you want to prove. It does *not* start by defining a Fourier cutoff and proving a generic localization inequality; instead it isolates the mean-field mechanism that can generate *geometric harmonic growth* when fitting a highly oscillatory target.

5.1 What must be proved (informal statement)

When the target is highly oscillatory, the trained network exhibits layerwise partial functions that are increasingly oscillatory and increasingly spike-like. The proof goal is to show that, in a sign-coherent regime, the mean-field gradient-flow update acts (after averaging over the gate) like a *sharpening operator* with gain strictly larger than 1, so that a characteristic spatial scale shrinks geometrically across depth, yielding geometric growth of accessible frequencies.

5.2 A minimal 1D mean-field sharpening model

Work on a 1D periodic domain $\mathbb{T}$. Let $u^{(\ell)}(t, \cdot)$ denote an effective *pre-ReLU* scalar field at depth ℓ (e.g. a dominant low-rank channel or a signed linear combination of channels). Define the *post-ReLU* field $v^{(\ell)}(t, x) = \max\{u^{(\ell)}(t, x), 0\}$. The empirical phenomenon is that if $u^{(\ell)}$ has many sign changes, then $v^{(\ell)}$ is supported on many short intervals, producing a collection of localized spikes.

The mean-field gradient flow updates $u^{(\ell)}$ through a gated expectation of the schematic form

$$\partial_t u^{(\ell)}(t, \cdot) \approx -\mathcal{K}_t^{(\ell)} \left(d(t, \cdot) \mathcal{G}_t^{(\ell)}[u^{(\ell)}(t, \cdot)] \right), \quad (35)$$

where $\mathcal{K}_t^{(\ell)}$ is a (layer-dependent) integral operator induced by the feature measure at depth ℓ , $d(t, \cdot)$ is the residual, and $\mathcal{G}_t^{(\ell)}$ is a *gate-induced* multiplication operator (coming from $\mathbb{1}\{S > 0\}$ and its correlations with the low-rank coordinates). The point is that $\mathcal{G}_t^{(\ell)}$ is *not* neutral: it biases updates toward the active set and can increase the effective slope/scale.

5.3 Half-space constants and why variance 2 matters

The basic half-space constants are explicit. If $G \sim \mathcal{N}(0, 2)$ then

$$\mathbb{E}|G| = \frac{2}{\sqrt{\pi}} > 1, \quad \mathbb{E}[G \mathbb{1}\{G > 0\}] = \frac{1}{\sqrt{\pi}}. \quad (36)$$

In the one-sparse / sign-coherent regimes used earlier, the conditional expectations over the gate reduce many vector-valued expressions to one-dimensional half-space averages. Equation (??) is exactly the place where the “variance 2” parameterization creates a gain constant > 1 .

5.4 A provable geometric harmonic-growth principle (template)

We now isolate a statement that captures the mechanism: *gate bias increases an effective slope parameter*, and slope growth implies shrinking spike widths, hence higher resolvable frequencies. To avoid introducing new Fourier notation, we quantify oscillation by the number of sign changes and by a characteristic spike width.

Definition 2 (Spike-width proxy). *For a nonnegative function $v : \mathbb{T} \rightarrow \mathbb{R}$ that is supported on a union of intervals, define its spike width proxy by*

$$h(v) \equiv \frac{\text{Leb}(\text{supp}(v))}{N(v)},$$

where $N(v)$ is the number of connected components of $\text{supp}(v)$ and Leb denotes Lebesgue measure on $\mathbb{T}$. When $v = u_+$ for an oscillatory u , $N(v)$ is proportional to the number of positive excursions, and $h(v)$ measures the typical width of a spike.

Assumption 2 (Mean-field gate sharpening across depth). *There exist deterministic constants $c_\ell > 1$ such that, on a regime where a single dominant channel controls the gate, the depth- ℓ to depth- $(\ell + 1)$ effective map satisfies the following implication: if $u^{(\ell)}(t, \cdot)$ has oscillation scale h_ℓ in the sense that $v^{(\ell)}(t, \cdot) = (u^{(\ell)}(t, \cdot))_+$ has spike width $h(v^{(\ell)}(t, \cdot)) \approx h_\ell$, then $v^{(\ell+1)}(t, \cdot)$ satisfies*

$$h(v^{(\ell+1)}(t, \cdot)) \lesssim \frac{1}{c_\ell} h(v^{(\ell)}(t, \cdot)), \quad \text{with} \quad c_\ell \approx \mathbb{E}|G|. \quad (37)$$

Lemma 14 (Geometric shrinkage implies geometric frequency growth). *Assume (??) holds along depths $1, \dots, L-1$ and that $c_\ell \geq c > 1$ uniformly. Then the typical spike width satisfies $h(v^{(\ell)}(t, \cdot)) \lesssim c^{-(\ell-1)} h(v^{(1)}(t, \cdot))$. In particular, the number of spikes satisfies*

$$N(v^{(\ell)}(t, \cdot)) \gtrsim c^{\ell-1} N(v^{(1)}(t, \cdot)) \quad \text{whenever } \text{Leb}(\text{supp}(v^{(\ell)})) \text{ stays bounded below.}$$

Thus, as depth increases, the representation necessarily contains structure at spatial scales of order $c^{-(\ell-1)}$, i.e. at frequencies of order $c^{\ell-1}$.

Proof deferred to appendix.

Remark 15 (What remains to turn this into a theorem for your MMNN mean-field PDE). *Lemma ?? makes explicit what must be proved: the core task is to justify Assumption ?? from the actual mean-field equations. Concretely, one needs to show that, when fitting an oscillatory target, the gate-induced operator $\mathcal{G}_t^{(\ell)}$ in (??) increases an effective slope parameter by a factor comparable to $\mathbb{E}|G|$ (with $G \sim \mathcal{N}(0, 2)$), and that this translates into the spike-width recursion (??). This is where the earlier sign-coherence / half-space bias lemmas are used: they reduce the gate statistics to one-dimensional half-space expectations, producing the gain constant in (??).*

5.5 Dictionary viewpoint and the 2^L oscillation pattern (why depth is the right lens)

Your observation can be rephrased in a way that is much closer to existing ReLU theory: *in 1D, deep ReLU networks represent splines whose number of breakpoints (hence oscillations) can grow exponentially with depth.* In the mean-field parameterization, the first layer already provides a universal dictionary; depth then controls how many times the input domain is *folded/cut* by gates, which is what creates many alternating spikes.

First layer as a universal dictionary. On 1D inputs, the ridge dictionary $x \mapsto (x - b)_+$ is universal for piecewise-linear splines. In the mean-field limit, choosing a signed measure over biases b and slopes is exactly choosing a (possibly very large) spline dictionary expansion. Thus the “rank- r components are basis functions” intuition is correct: at each layer, a small number of channel mixtures selects and reweights a huge underlying dictionary.

Depth creates exponentially many breakpoints. The following is a standard expressivity fact in 1D: with ReLU, the function at each hidden unit is piecewise-linear, and each new layer can multiply the number of linear pieces. This yields an exponential-in-depth upper bound, and (by explicit constructions) an exponential-in-depth *lower* bound is achievable.

Definition 3 (Oscillation count). *For a continuous piecewise-linear function $f : \mathbb{T} \rightarrow \mathbb{R}$, let $\text{osc}(f)$ denote the number of maximal intervals on which f is affine (equivalently, one plus the number of breakpoints). This quantity controls the number of sign changes and the number of spikes after thresholding.*

Lemma 15 (1D deep ReLU networks are splines with exponentially many pieces). *Consider a standard fully-connected ReLU network on 1D input $x \in \mathbb{T}$, of width m and depth L (i.e. L hidden layers), with scalar output. Then the realized function is continuous piecewise-linear and satisfies*

$$\text{osc}(f) \leq C m^L$$

for a universal constant C (e.g. one may take $C = 2$).

Proof deferred to appendix.

Lemma 16 (Exponential oscillations are achievable (sawtooth / folding construction)). *There exist explicit ReLU networks on $x \in [0, 1]$ whose output has $\text{osc}(f) \geq 2^L$ after L repeated “folding” steps. Moreover, such foldings can be implemented with constant width by using a small fixed ReLU gadget and stacking it L times (equivalently, with depth proportional to L).*

Proof deferred to appendix.

Remark 16 (Why this matches your “one block $\approx$ two hidden layers” convention). *Many explicit folding constructions use a constant-size ReLU gadget whose implementation may take two affine-ReLU stages. If your notion of a “layer” corresponds to such a block (from input to the r low-rank components, with gating/mixing inside), then the empirical pattern “depth L produces about 2^L oscillations” is consistent with Lemma ?? . For larger rank/width, the natural scaling is $\text{osc}(f)$ of order r^L .*

What this reduces the training proof to. The expressivity lemmas above do *not* say that gradient flow will realize the worst-case oscillation count. But they strongly suggest the right invariant/proxy to analyze is not a Fourier cutoff: it is a *knot/oscillation creation functional*, e.g. $\text{osc}(u^{(\ell)})$ or a smooth surrogate such as $\text{TV}(\partial_x u^{(\ell)})$. The mean-field task is then to show that, when the residual alternates sign at scale λ^{-1} , the gated dynamics create additional breakpoints at roughly that scale and that the gain constant in (??) makes this process stable across depth.